

```
In [90]: #import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
import sklearn
import seaborn as sns
!pip install haversine
```

Requirement already satisfied: haversine in c:\users\dell\appdata\local\programs\python\python310\lib\site-packages (2.8.0)

```
In [91]: #import data
data = pd.read_csv("uber.csv")
```

```
In [92]: data
```

Out[92]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.750000
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.750000
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.750000
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.750000
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.750000
...	...	...	...	...	...	...
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	40.750000
199996	16382965	2014-03-14 01:09:00.0000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	40.750000
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	40.750000
199998	20259894	2015-05-20 14:56:25.0000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	40.750000
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	40.750000

200000 rows × 7 columns



```
In [93]: #Create a data copy
df = data.copy()
```

```
In [94]: #Print data
df=df.head(8000)
```

```
In [95]: #Get Info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8000 entries, 0 to 7999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            8000 non-null   int64
1   key                   8000 non-null   object
2   fare_amount           8000 non-null   float64
3   pickup_datetime       8000 non-null   object
4   pickup_longitude      8000 non-null   float64
5   pickup_latitude       8000 non-null   float64
6   dropoff_longitude     8000 non-null   float64
7   dropoff_latitude      8000 non-null   float64
8   passenger_count       8000 non-null   int64
dtypes: float64(5), int64(2), object(2)
memory usage: 562.6+ KB
```

```
In [96]: #pickup_datetime is not in required data format
df["pickup_datetime"] = pd.to_datetime(df["pickup_datetime"])
```

```
In [97]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8000 entries, 0 to 7999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            8000 non-null   int64
1   key                   8000 non-null   object
2   fare_amount           8000 non-null   float64
3   pickup_datetime       8000 non-null   datetime64[ns, UTC]
4   pickup_longitude      8000 non-null   float64
5   pickup_latitude       8000 non-null   float64
6   dropoff_longitude     8000 non-null   float64
7   dropoff_latitude      8000 non-null   float64
8   passenger_count       8000 non-null   int64
dtypes: datetime64[ns, UTC](1), float64(5), int64(2), object(1)
memory usage: 562.6+ KB
```

In [98]: *#Statistics of data*  
df.describe()

Out[98]:

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	8.000000e+03	8000.000000	8000.000000	8000.000000	8000.000000	8000.000000
mean	2.769486e+07	11.442829	-72.674559	39.981164	-72.635342	39.981164
std	1.595445e+07	10.467626	12.665719	6.006576	10.007892	6.006576
min	4.800000e+02	2.500000	-74.016667	-74.009697	-75.350437	-73.981164
25%	1.402794e+07	6.000000	-73.992066	40.735101	-73.991471	40.735101
50%	2.764842e+07	8.500000	-73.981504	40.752477	-73.979967	40.752477
75%	4.127959e+07	12.500000	-73.967069	40.766865	-73.963482	40.766865
max	5.542169e+07	350.000000	40.770667	41.366138	40.761672	41.366138

In [99]: *#Number of missing values*  
df.isnull().sum()

Out[99]:

Unnamed: 0	0
key	0
fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	0
dropoff_latitude	0
passenger_count	0
dtype:	int64

In [100]:  
numeric\_df = df.select\_dtypes(include=[np.number])  
correlation\_matrix = numeric\_df.corr()  
correlation\_matrix

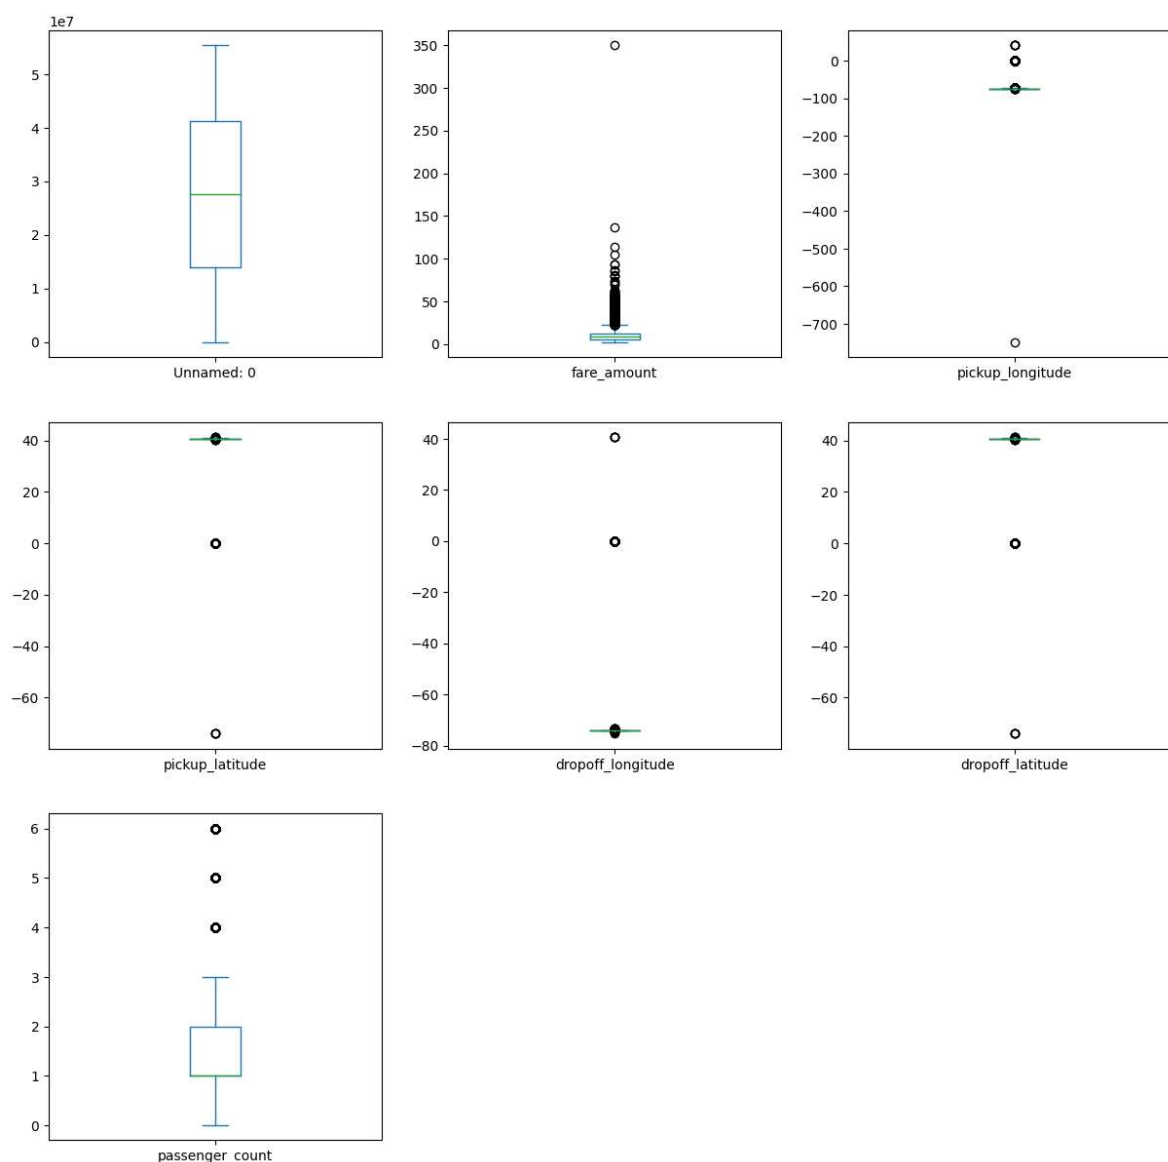
Out[100]:

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
Unnamed: 0	1.000000	0.005721	0.013433	-0.025790	0.026929	-0.025790
fare_amount	0.005721	1.000000	0.037397	-0.037727	0.047430	-0.037727
pickup_longitude	0.013433	0.037397	1.000000	-0.773625	0.758229	-0.773625
pickup_latitude	-0.025790	-0.037727	-0.773625	1.000000	-0.934771	1.000000
dropoff_longitude	0.026929	0.047430	0.758229	-0.934771	1.000000	-0.934771
dropoff_latitude	-0.025790	-0.037727	-0.773625	1.000000	-0.934771	1.000000
passenger_count	0.009653	-0.012580	0.010977	-0.013181	0.012631	-0.013181

```
In [101]: #Drop the rows with missing values
df.dropna(inplace=True)
```

```
In [102]: df.plot(kind="box", subplots=True, layout = (4,3),figsize=(15,20))
```

```
Out[102]: Unnamed: 0      AxesSubplot(0.125,0.712609;0.227941x0.167391)
fare_amount      AxesSubplot(0.398529,0.712609;0.227941x0.167391)
pickup_longitude AxesSubplot(0.672059,0.712609;0.227941x0.167391)
pickup_latitude  AxesSubplot(0.125,0.511739;0.227941x0.167391)
dropoff_longitude AxesSubplot(0.398529,0.511739;0.227941x0.167391)
dropoff_latitude AxesSubplot(0.672059,0.511739;0.227941x0.167391)
passenger_count  AxesSubplot(0.125,0.31087;0.227941x0.167391)
dtype: object
```



```

In [103]: def remove_outlier(df1, col):
            if df1[col].dtype in [int, float]:
                Q1= df1[col].quantile(0.25)
                Q2=df1[col].quantile(0.50)

                Q3=df1[col].quantile(0.75)

                IQR = Q3-Q1

                lower_whisker = Q1-1.5*IQR

                upper_whisker= Q3+1.5*IQR
                print("col=",col, "Q1=",Q1,"Q2=", Q2,"Q3=",Q3)
                df1[col] = np.clip(df1[col], lower_whisker, upper_whisker)
            else:
                print(f"Column {col} is not numeric and cannot calculate quantiles.")
            return df1

def treat_outliers_all(df1, col_list):

    print("col_list",col_list)

    for c in col_list:

        df1 = remove_outlier(df1, c)

    return df1
df= treat_outliers_all(df, df.columns)

```

```

col_list Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
               'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
               'dropoff_latitude', 'passenger_count'],
              dtype='object')

```

Column Unnamed: 0 is not numeric and cannot calculate quantiles.

Column key is not numeric and cannot calculate quantiles.

col= fare\_amount Q1= 6.0 Q2= 8.5 Q3= 12.5

Column pickup\_datetime is not numeric and cannot calculate quantiles.

col= pickup\_longitude Q1= -73.992066 Q2= -73.981504 Q3= -73.96706922595214

col= pickup\_latitude Q1= 40.73510125 Q2= 40.7524765 Q3= 40.766865


col= dropoff\_longitude Q1= -73.99147124999999 Q2= -73.97996679028321 Q3= -73.9634815

col= dropoff\_latitude Q1= 40.733695 Q2= 40.753132 Q3= 40.768231

Column passenger\_count is not numeric and cannot calculate quantiles.

```
In [104]: df.plot(kind="box",subplots= True , layout=(7,3), figsize=(15,20))
```

```
Out[104]: Unnamed: 0      AxesSubplot(0.125,0.786098;0.227941x0.0939024)
fare_amount      AxesSubplot(0.398529,0.786098;0.227941x0.0939024)
pickup_longitude AxesSubplot(0.672059,0.786098;0.227941x0.0939024)
pickup_latitude  AxesSubplot(0.125,0.673415;0.227941x0.0939024)
dropoff_longitude AxesSubplot(0.398529,0.673415;0.227941x0.0939024)
dropoff_latitude AxesSubplot(0.672059,0.673415;0.227941x0.0939024)
passenger_count  AxesSubplot(0.125,0.560732;0.227941x0.0939024)
dtype: object
```



```
In [105]: #Check the missing values now
df.isnull().sum()
```

```
Out[105]: Unnamed: 0      0
key      0
fare_amount      0
pickup_datetime      0
pickup_longitude      0
pickup_latitude      0
dropoff_longitude      0
dropoff_latitude      0
passenger_count      0
dtype: int64
```

```
In [106]: #Time to apply learning models
```

```
from sklearn.model_selection import train_test_split
```

```
In [107]: #Take x as predictor variable  
x = df.drop("fare_amount", axis = 1)  
#And y as target variable  
y = df['fare_amount']
```

```
In [108]: #Necessary to apply model  
x['pickup_datetime'] = pd.to_numeric(pd.to_datetime(x['pickup_datetime']))  
x = x.loc[:, x.columns.str.contains('^Unnamed')]
```

```
In [109]: import haversine as hs

# Define a function to calculate the haversine distance
def haversine_distance(row):
    loc1 = (row['pickup_latitude'], row['pickup_longitude'])
    loc2 = (row['dropoff_latitude'], row['dropoff_longitude'])
    return hs.haversine(loc1, loc2)

# Apply the function to each row and create a new column
df['dist_travel_km'] = df.apply(haversine_distance, axis=1)

# Print the distances and display the DataFrame
print(df['dist_travel_km'])
print(df.head())
```

```
0      1.683325
1      2.457593
2      5.036384
3      1.661686
4      4.131933
```

```
...
7995   11.792129
7996    1.714758
7997    0.990416
7998    5.714015
7999    0.920174
```

Name: dist\_travel\_km, Length: 8000, dtype: float64

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.0000003	7.5	
1	27835199	2009-07-17 20:04:56.0000002	7.7	
2	44984355	2009-08-24 21:45:00.00000061	12.9	
3	25894730	2009-06-26 08:22:21.0000001	5.3	
4	17610152	2014-08-28 17:47:00.000000188	16.0	

	pickup_datetime	pickup_longitude	pickup_latitude	\
0	2015-05-07 19:52:06+00:00	-73.999817	40.738354	
1	2009-07-17 20:04:56+00:00	-73.994355	40.728225	
2	2009-08-24 21:45:00+00:00	-74.005043	40.740770	
3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	
4	2014-08-28 17:47:00+00:00	-73.929574	40.744085	

	dropoff_longitude	dropoff_latitude	passenger_count	dist_travel_km
0	-73.999512	40.723217	1	1.683325
1	-73.994710	40.750325	1	2.457593
2	-73.962565	40.772647	1	5.036384
3	-73.965316	40.803349	3	1.661686
4	-73.973082	40.761247	5	4.131933

```
In [110]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, ran
```

```
In [111]: from sklearn.linear_model import LinearRegression
```



```
In [112]: correlation_matrix = df.corr()

# Print the correlation matrix
print(correlation_matrix)
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
```

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude
Unnamed: 0	1.000000	0.005938	0.008992	-0.004565
fare_amount	0.005938	1.000000	0.165482	-0.115839
pickup_longitude	0.008992	0.165482	1.000000	0.272431
pickup_latitude	-0.004565	-0.115839	0.272431	1.000000
dropoff_longitude	0.023527	0.223993	0.415318	0.059101
dropoff_latitude	-0.007352	-0.134131	0.090473	0.518721
passenger_count	0.009653	-0.007241	-0.009506	-0.010992
dist_travel_km	-0.004231	0.794487	0.058659	-0.079959


	dropoff_longitude	dropoff_latitude	passenger_count	\
Unnamed: 0	0.023527	-0.007352	0.009653	
fare_amount	0.223993	-0.134131	-0.007241	
pickup_longitude	0.415318	0.090473	-0.009506	
pickup_latitude	0.059101	0.518721	-0.010992	
dropoff_longitude	1.000000	0.260177	-0.010252	
dropoff_latitude	0.260177	1.000000	-0.018341	
passenger_count	-0.010252	-0.018341	1.000000	
dist_travel_km	0.157602	-0.072390	-0.001882	

	dist_travel_km
Unnamed: 0	-0.004231
fare_amount	0.794487
pickup_longitude	0.058659
pickup_latitude	-0.079959
dropoff_longitude	0.157602
dropoff_latitude	-0.072390
passenger_count	-0.001882
dist_travel_km	1.000000

C:\Users\Dell\AppData\Local\Temp\ipykernel\_12836\1932690646.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
correlation_matrix = df.corr()
```

```
Out[112]: <AxesSubplot: >
```



```
In [113]: lrmodel = LinearRegression()
lrmodel.fit(x_train, y_train)
```

```
Out[113]: ▾ LinearRegression
LinearRegression()
```

```
In [127]: print(lrmodel.intercept_)
print(lrmodel.coef_)
```

```
10.13702658872479
[1.80835081e-10]
```

```
In [114]: #Prediction
predict = lrmodel.predict(x_test)
```

```
In [133]: comparison=pd.DataFrame({"Actual_Price":y_test,"Predicted_Price":predict})  
print(comparison.reset_index().drop(["index"],axis=1))  
sns.heatmap(comparison.corr())
```

	Actual_Price	Predicted_Price
0	7.0	10.143457
1	11.3	10.144738
2	4.5	10.139688
3	6.5	10.142855
4	5.5	10.143302
...	...	...
1595	7.0	10.140242
1596	9.5	10.146518
1597	9.7	10.140564
1598	4.5	10.146202
1599	8.9	10.143957

[1600 rows x 2 columns]

Out[133]: <AxesSubplot: >



```
In [115]: #Check Error
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.metrics import r2_score
predict = lrmodel.predict(x_test)
lr_r2 = r2_score(y_test, predict)
lrmodelrmse = np.sqrt(mean_squared_error(predict, y_test))
mse = mean_squared_error(y_test, predict)
print(f"Mean Squared Error (MSE): {mse:.2f}")
mae = mean_absolute_error(y_test, predict)
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print("Linear Regression R2:", lr_r2)
print("RMSE error for the model is ", lrmodelrmse)
```

Mean Squared Error (MSE): 28.68  
Mean Absolute Error (MAE): 4.33  
Linear Regression R2: -0.00204461954986912  
RMSE error for the model is 5.355275698890742

```
In [116]: #Let's Apply Random Forest Regressor
from sklearn.ensemble import RandomForestRegressor
rfrmodel = RandomForestRegressor(n_estimators = 100, random_state = 101)
```

```
In [117]: #Fit the Forest
rfrmodel.fit(x_train, y_train)
rfrmodel_pred = rfrmodel.predict(x_test)
```

```
In [118]: #Errors for the forest

rf_r2 = r2_score(y_test, rfrmodel_pred)
rfrmodel_rmse = np.sqrt(mean_squared_error(rfrmodel_pred, y_test))
print("Random Forest Regression R2:", rf_r2)
print("RMSE value for Random Forest is:", rfrmodel_rmse)
mse = mean_squared_error(y_test, rfrmodel_pred)
print(f"Mean Squared Error (MSE): {mse:.2f}")
mae = mean_absolute_error(y_test, rfrmodel_pred)
print(f"Mean Absolute Error (MAE): {mae:.2f}")
```

Random Forest Regression R2: -0.5489964251584396  
RMSE value for Random Forest is: 6.658302336824492  
Mean Squared Error (MSE): 44.33  
Mean Absolute Error (MAE): 5.18

```
In [119]: print(rfrmodel_pred)
print(predict)
```

```
[ 9.156 10.074  9.01  ...  7.467  8.4095 13.493 ]
[10.14345698 10.14473803 10.13968762 ... 10.14056434 10.1462023
 10.14395714]
```

